

# High Performance Stock Trading System

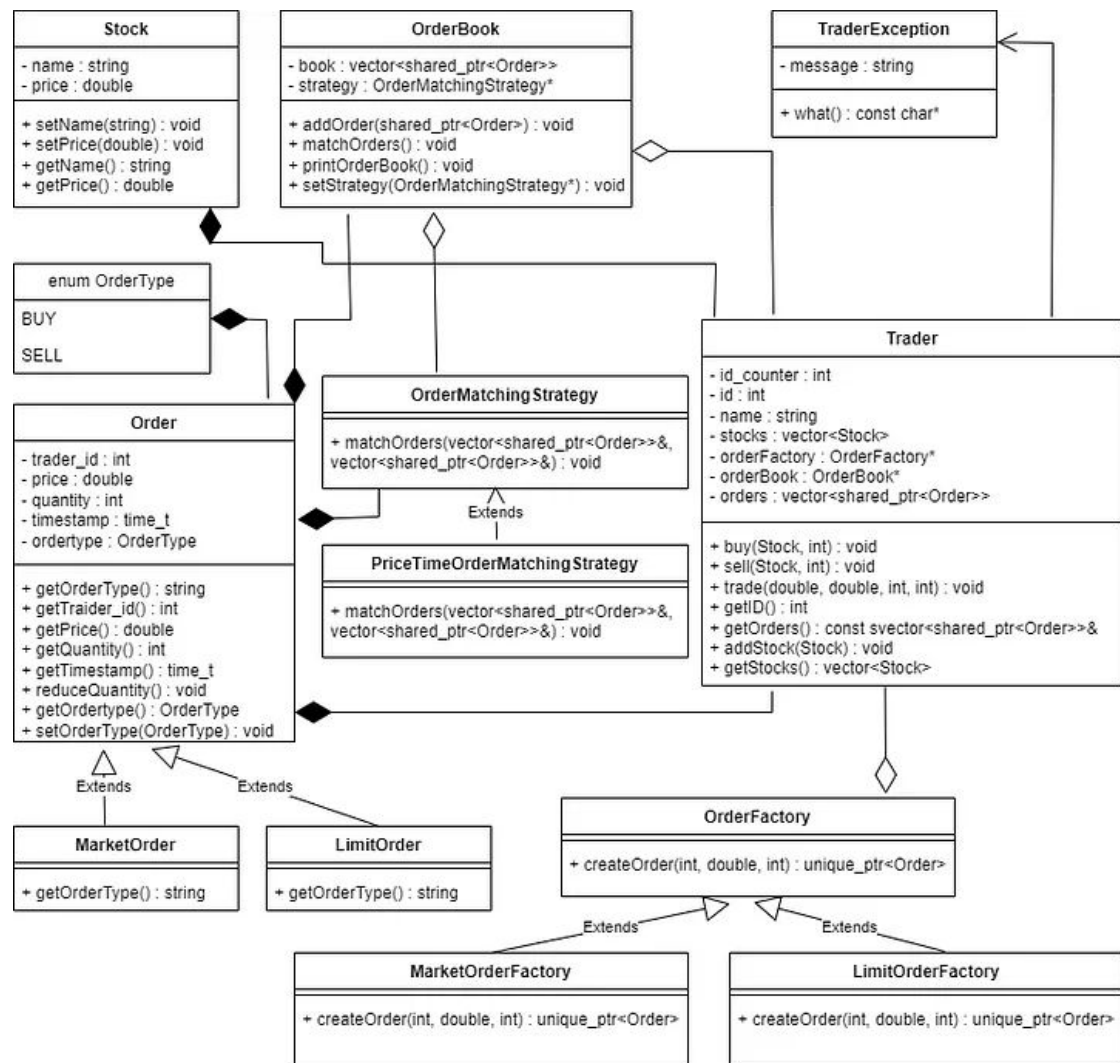
CS205 Course Project Proposal

Group 8



## Overview

- Simulation of realistic stock trading platform
- Core functionalities:
  - Receive, put, and call orders
  - Access stock status
  - Update orderbook according to the received requests
  - Match buy and sell orders according to order types (e.g. market/limit) and specified strategies
  - Execute matched orders and update orderbook and stock status accordingly



# Simulated Trading Data

- Data Generation

- We wrote a script that can generate large datasets of traders and stocks with varying quantities and complexities
- Ensures a high degree of randomness, essential for simulating realistic financial market behaviors.
- **Traders**
  - unique names & randomized financial thresholds and quantities, to simulate diverse market participant behaviors
- **Stocks**
  - unique tickers and prices, using uniform distribution to model market price variations realistically.

- avoids the complexities and unpredictabilities of real-time financial data, such as network latencies and data inconsistencies
- prioritizes the creation of a controlled, predictable environment for reliable HPC benchmarking and testing
- stocks number from **5k to 500k**, so size from **2\*5K to 2\*500K**, and traders number from **1K to 14M**, so size from **5\*1K to 5\*14M**

```
/*
traders.dat
...
name, buyThreshold, sellThreshold, buyQuantity, sellQuantity
...
types:
string, double, double, int, int
*/

/*
stocks.dat
...
name, price
...
types:
string, double
*/
```

# Limitations of Sequential Implementation



Suppose we have  $N$  stocks,  $T$  traders, and  $M$  orders in the system:

- Operations done by each trader:
  - Retrieve a shared list of stocks:  $O(N * T)$
  - Put orders into orderbook:  $O(N * T)$
- Match and execute orders:  $O(M^2)$
- Update the orderbook, the stocks, and the user accounts:  $O(M+N+T)$

The upper bound of  $M$  is also  $O(N * T)$

Performance would degrade when these numbers are large.

# Parallelization (Part 1)



- Order Creation Process
  - User access to current stock status can also be performed in parallel
  - Putting and calling orders can be read in parallel and sorted in the orderbook
- Order-Matching Process
  - Order-Level Parallelization (**OpenMP, MPI**)
    - Maybe partially parallelized depending on matching strategy
    - Unordered requests can be segmented into chunks by timestamp and handled separately by each thread
    - Final ordering w.r.t timestamp chunks needs to be done in serial manner

# Parallelization (Part 2)



- Execution (Update) Process
  - Stock-Level Parallelization (**MPI**)
    - Potentially fully parallelized
    - Different stocks are independent, orders w.r.t each stock type can be totally separated and updated in parallel
  - User-Level Parallelization (*bonus*: if bandwidth allowed) (**MPI**)
    - Potentially fully parallelized
    - Keep track of the income and spend amount of each user (i.e., user accounts)
    - Putting and calling results can be stored in random order
    - Only needs 1 final integration step to update user account information at the end of a certain time period (i.e., one trading day).